


Business intelligence

Unit 3 – Data exploitation. Query languages
and visualization

S3-1 – OLAP

OLAP tools are used in BI for **analyzing multidimensional data** from multiple perspectives. OLAP tools provide the user with a **multidimensional view** of data (multidimensional schema) for each activity that is being analyzed. The user formulates queries to the OLAP tool selecting multidimensional attributes of this scheme **without knowing the internal structure** (physical schema) of the data warehouse. The tool generates a corresponding **OLAP query** and sends it to the query management system (e.g. by means a SQL SELECT statement).




Multidimensional Analysis: OLAP enables **slicing, dicing, and drilling down** into data for additional insights.

Complex Calculations: OLAP tools support advanced calculations and metrics to evaluate business performance.

Fast Query Performance: MOLAP provides pre-aggregation, while HOLAP offers hybrid performance improvements.


Data Consistency: Ensures accurate, consistent analysis across BI applications.

Query resolution procedure: Build the query → Extract → aggregated data → Visualize results → Analyze

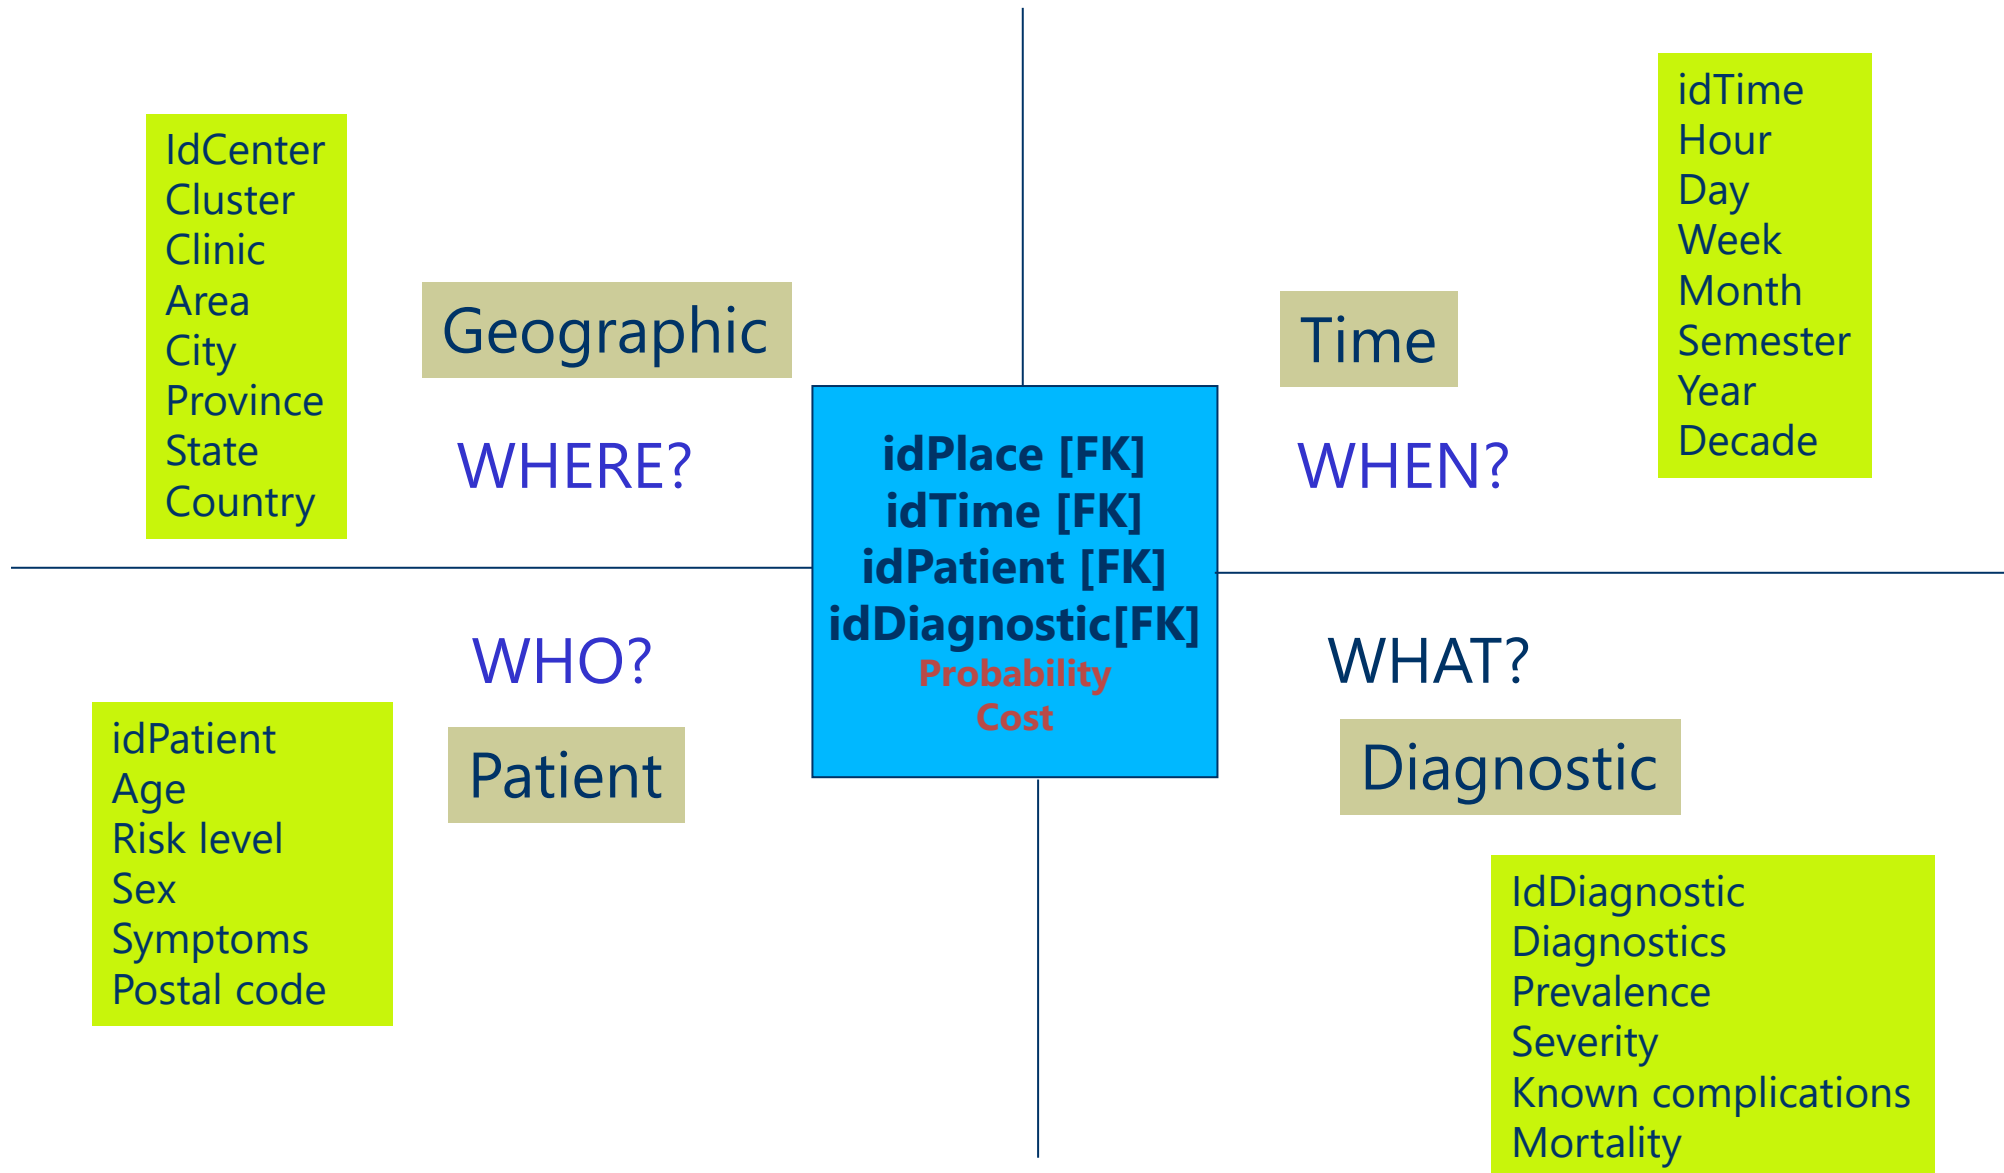



An **OLAP query** consists of

- Retrieve measures or indicators
- About the facts
- parametrized by attributes in the dimensions
- Constrained by conditions imposed on the dimensions



Eg: What is the total cost per diagnostic with low mortality rate in the last year for each province and sex?





Fact tables

Diagnostic	Sex	Total
D1	M	100
D1	F	200
D2	M	150
D2	F	75



2D view

Diag\Sex	M	F
D1	100	200
D2	150	75

Fact table

Diagnostic	Sex	Province	Total
D1	M	P1	100
D1	F	P1	200
D1	M	P1	100
D1	F	P1	200
D2	M	P2	150
D2	F	P2	75
D2	M	P2	150
D2	F	P2	75

3D view


The 3D view represents the fact table data in a cube structure. The vertical axis is labeled 'Province' with values P1 and P2. The horizontal axis is labeled 'Diagnostic' with values D1 and D2. The depth axis is labeled 'Sex' with values M and F. The data values are shown in the cells of the cube.

P2	Diag\Sex	M	F
	D1	100	200
P1	D2	150	75
	Diag\Sex	M	F
P1	D1	100	200
	D2	150	75

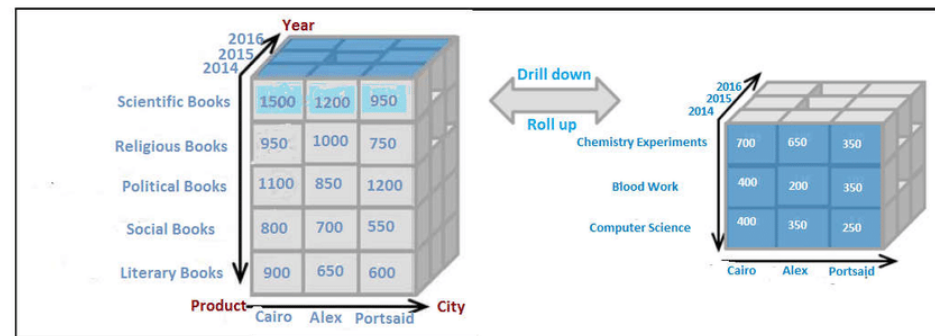
The interesting thing is **NOT ONLY** to be able to query, in a way, something you can do with selections, projections, concatenation and traditional groupings.

What is really interesting OLAP tools are its refinement operators for handling queries.

DRILL
ROLL
SLICE & DICE
PIVOT
ROLLUP
CUBE



Diagnostic	Sex	Province	Total
D1	M	P1	100
D1	F	P1	200
D1	M	P1	100
D1	F	P1	200
D2	M	P2	150
D2	F	P2	75
D2	M	P2	150
D2	F	P2	75



Diagnostic	Sex	Total
D1	M	200
D1	F	400
D2	M	300
D2	F	150


roll →

← drill

Aggregate (consolidate) and **disintegrate** (division):

- aggregation (**roll**): delete a grouping criterion in the analysis, aggregating the current groups. The granularity of one or more dimensions is aggregated.
- disintegrate (**drill**): enter a new grouping criterion in the analysis, breaking existing groups.

Aggregation in SQL: sum, count, max, min, average.

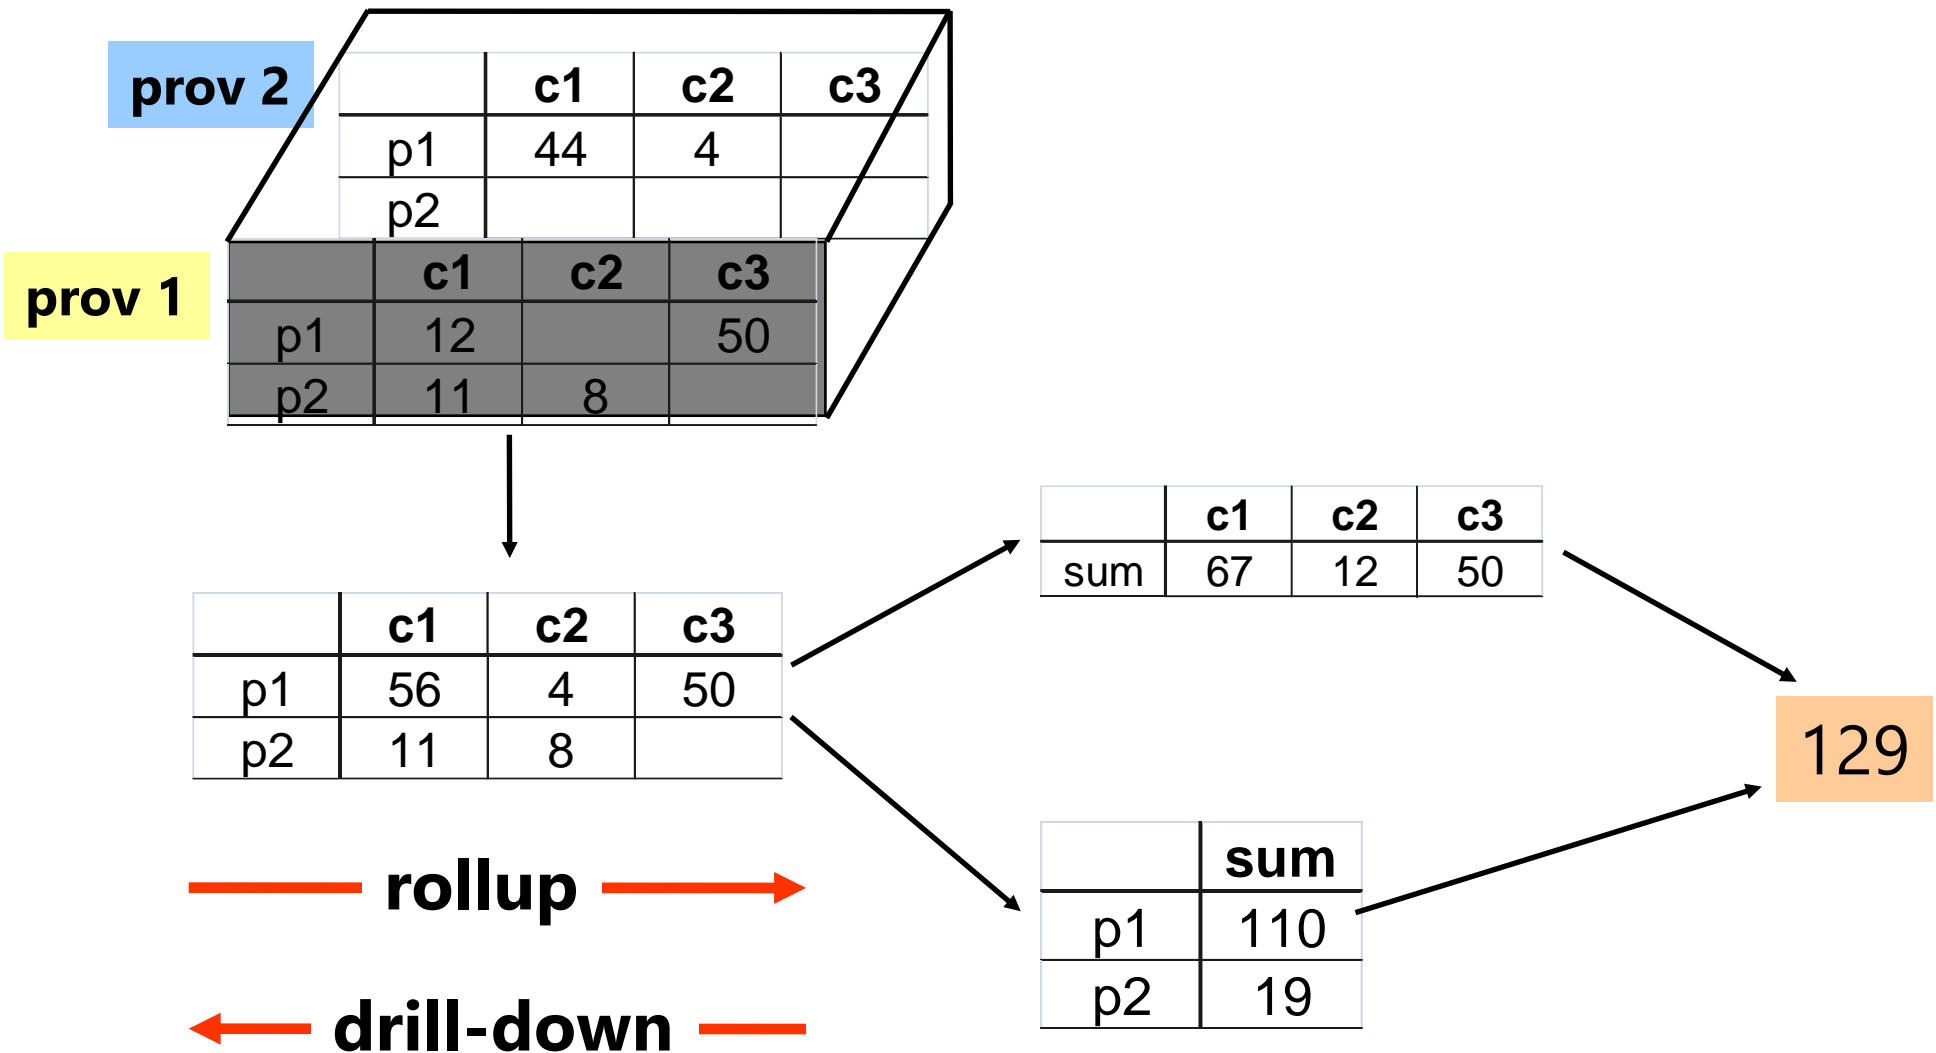


DRILL (ROLL) can be done on:

- **attributes** of one dimension on which a hierarchy has been defined:
 - DRILL-DOWN: upper to lower aggregation level
 - departament – category - product (Product)
 - year – semester – month – day (Time)
 - ROLL-UP: lower to upper aggregation level.

Other “drill”:

- **DRILL-ACROSS:** join several fact tables.
- **DRILL-THROUGH:** Use SQL to explore up to the relational back-end tables.




SLICE & DICE: select and project

SLICE: Filter a specific dimension by selecting a single value.

DICE: Filter data by applying multiple conditions across dimensions.

PIVOT: Rotate, reorientate the data view to analyze different perspectives.

Store	Product	Sales
A	TV	2
A	TV	4
B	TV	6
B	DVD	8



Store	Avg(Sales) for TV	Avg(Sales) for DVD
A	3	(Empty)
B	6	8

Diagnóstico	Sexo	Provincia	Total	Núm
D1	H	P1	100	6
D1	M	P1	200	5
D1	H	P2	100	6
D1	M	P2	200	11
D2	H	P1	150	7
D2	M	P1	75	7
D2	H	P2	150	2
D2	M	P2	75	1

**Slice by removing
Prov dimension**



Diagnostic	Sex	Total
D1	M	100
D1	F	200
D2	M	150
D2	F	70

```
SELECT diagnostic, sex,
SUM(total)
FROM table
WHERE province = 'P1'
GROUP BY diagnostic, sex;
```

	Diagnóstico	Sexo	Total
P1	D1	H	100
	D1	M	200
	D2	H	150
	D2	M	75
P2	D1	H	100
	D1	M	200
	D2	H	150
	D2	M	75

Pivot
→

	Diagnóstico	Prov incia	Total
I	D1	P1	100
	D1	P2	100
	D2	P1	150
	D2	P2	150
M	D1	P1	200
	D1	P2	200
	D2	P1	75
	D2	P2	75

Example

```
CREATE EXTENSION IF NOT EXISTS tablefunc;
```

```
CREATE TABLE to_pivot (  
  ID serial,  
  Name TEXT, -- Name student  
  Course TEXT, -- Course  
  Grade INT,  
  primary KEY(ID)  
);
```

```
INSERT INTO to_pivot(Name, Course, Grade) VALUES
```

```
( 'Pepe', 'BDII', 9),  
( 'Jose', 'BDII', 7),  
( 'Pepe', 'BI', 8),  
( 'Jose', 'BI', 5);
```

	123 id	A-Z name	A-Z course	123 grade
1	1	Pepe	BDII	9
2	2	Jose	BDII	7
3	3	Pepe	BI	8
4	4	Jose	BI	5

	A-Z name	123 BI	123 BDII
1	Jose	7	5
2	Pepe	9	8

```
SELECT * FROM to_pivot;
```

```
SELECT * FROM crosstab ('Select Name, Course, Grade from to_pivot order by  
1, 2') as Pivoted (Name text, "BI" INT, "BDII" INT);
```


SQL aggregation

- sum(), count(), avg(), min(), max()

Basic idea:

- Combine values in one column
- Into only one value


Syntax:

```
SELECT sum(cost) FROM diagnostic;
```

DISTINCT

- Allows the aggregation only of different values

```
SELECT COUNT(DISTINCT cost) FROM diagnostic;
```



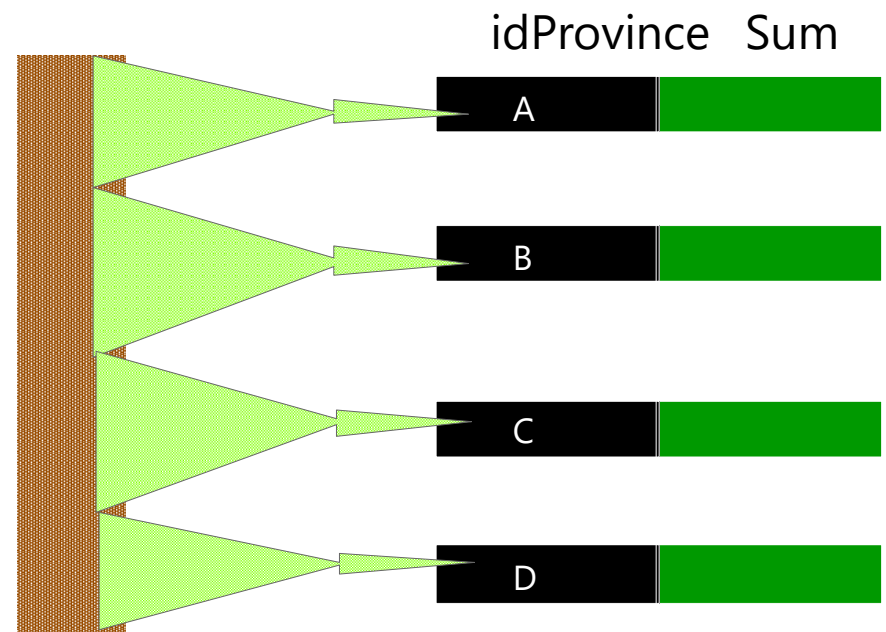
GROUP BY: Groups rows that have same values in specified columns. Aggregation functions are usually associated to it: **Select** idprov, **sum**(b) **from** table **group by** idprov

GROUP BY + HAVING

Aggregating in subgroups of the table that fulfill some condition applied after the grouping.

Syntax


```
SELECT idProvinc, sum(cost)
FROM diagnostic
GROUP BY idProvinc
HAVING population > 2000;
```



Limitations without them

- Useful aggregations are difficult to calculate
 - Data cube
 - Complex: median, variance
 - Moving average
 - Rankings
- Marginals or crosstabs
 - SQL requires additional functionality
- Include sum and partial sums
 - drill-down & roll-up

	M	T	W	T	F	S	S.
AIR							
HOTEL							
FOOD							
MISC							



ROLLUP: performs the aggregation for the set of prefix of the attributes given

Example:

```
SELECT item-name, color, size, SUM(number)  
FROM sales
```

```
GROUP BY ROLLUP(item-name, color, size)
```

Calculates **SUM for the n+1 prefixes:**

{ (item-name, color, size), (item-name, color), (item-name), () }

Very useful for aggregating in hierarchies defined on dimensions

It can be done in SQL without OLAP extensions, but very inefficiently (multiple GROUP BY and UNION operations).

To improve efficiency: calculate the higher level aggregations using partial results of the more detailed levels

CUBE: generalization of GROUP BY to n-dimensions.

- Calculates the aggregation function for all the subsets of the attributes given instead for only the prefixes (ROLLUP)
- Example:

```
SELECT item-name, color, size, SUM(number)
FROM sales
GROUP BY CUBE (item-name, color, size)
```

- Calculates the aggregate for the set of 2^n combinations:
- { (item-name, color, size),
 (item-name, color), (item-name, size), (color, size),
 (item-name), (color), (size),
 () }
- For each combination, the result is null for attributes that are not present in the combination.

SQL:1999 uses NULL for representing both aggregated rows (ALL) and “usual” null (missing values).

When we have an OLAP query, how to know?

- In order to distinguish them we can use the **GROUPING** function that applied to an attribute
 - Returns 1 if NULL represents ALL
 - Returns 0 otherwise
 - Combined with DECODE (or CASE) we can return the desired value
- ```
SELECT DECODE(GROUPING(Year), 1, 'Total',
Year) AS Year, DECODE(GROUPING(Region), 1, 'Total', Region) AS
Region, SUM(SalesAmount) AS TotalSales FROM Sales GROUP BY CUBE
(Year, Region);
```

**GROUPING SETS** allows us to specify multiple groupings in a single query.

We can define the subsets of columns for grouping.

No need to have separate queries or UNION ALL.

Provides better efficiency.

```
select Name, Course, AVG(Grade) from
to_pivot group by
grouping sets
((Name, Course), (Course), ());
```

|   | A-Z name | A-Z course | 123 avg |
|---|----------|------------|---------|
| 1 | [NULL]   | [NULL]     | 7.25    |
| 2 | Pepe     | BI         | 8       |
| 3 | Jose     | BI         | 5       |
| 4 | Jose     | BDII       | 7       |
| 5 | Pepe     | BDII       | 9       |
| 6 | [NULL]   | BDII       | 8       |
| 7 | [NULL]   | BI         | 6.5     |

**WINDOW** clause defines **ordered** and **overlapping** groups of rows to calculate aggregates based on a defined "window", while retaining the original rows.

**GROUP BY** clause defines disjoint partitions of tuples in a sorted table, then calculates aggregates on those partitions, and generates a tuple with the result of the aggregate for each partition. It eliminates rows-level granularity

- Example: "For each day, we want the average cost of obtaining diagnoses from the previous day, the current and the next, and cumulatively in the last 7 days":

```
SELECT date, sum(cost) OVER (order by date ROWS BETWEEN 1
preceding and 1 following), sum(cost) OVER (order by date ROWS
BETWEEN 7 preceding and CURRENT ROW)) FROM diagnostics;
```



## Syntax:

- SELECT attribute\_list\_1, + Aggregated\_function **OVER** W as windowName
- FROM table\_list
- WHERE constraints

## **WINDOW** W AS (

- **PARTITION BY** attribute\_list\_2
- **ORDER BY** attribute\_list\_3
- **frame declaration**)

Frame declaration is opcional. By default, RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

## Execution:

- FROM, WHERE, GROUP and HAVING generate an **intermediate** table.
- PARTITION: each partition contains tuples with the same values in the attributes given in attribute\_list\_2
- ORDER BY: rows in each partition are sorted according to the values of the attributes in attribute\_list\_3
- SELECT the tuples under the constraints established in the frame declaration
  - RANGE: logical conditions (ie: 5 days)
  - ROWS: in rows (ie: 5 preceding rows)

## Frame examples:

- between rows unbounded preceding and current row
- rows unbounded preceding
- range between 10 preceding and current row
- range interval 10 day preceding
- range between interval 1 month preceding and interval 1 month following

**Default frame:** If the frame is not specified, all preceding and current rows are considered in the partition

- RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

**RANK assigns** to every tuple a rank based in some sorting of some attribute

- Example: given a cost-province relation rank each province by its cost.

```
SELECT province, rank() over (order by coste desc) as
provrank FROM diagnostic
```

- Afterwards, the result can be sorted by that field

```
SELECT province, rank() over (order by coste desc) as
provrank FROM diagnostic order by provrank
```

**RANK** allow gaps if there are 2 values with the same ranking.

- Example: if the 1<sup>st</sup> and 2<sup>nd</sup> classified have the same cost, then both will be assigned rank 1, and the next row will have rank 3
- DENSE\_RANK does not allow gaps, so the next row will have rank 2

- RANK over partitions:
  - "Rank the community and provinces by their cost"

```
SELECT province, comunity,
rank () over (partition by comunity order by cost
desc) as prov-comunity-rank
FROM diagnostic
ORDER BY comunity, prov-comunity-rank
```

si particionamos por provincia, daría que todas las provincias tendrían ranking 1

- Several RANK can be included in the same query.

## Other rank functions

- **percent\_rank**: it displays each row as a percentage of all the other rows up to 100% in a rank.  $= (\text{rank} - 1) / (\text{Number\_rows\_partition} - 1)$ .

```
SELECT Student, Score,
PERCENT_RANK() OVER
(ORDER BY Score DESC) AS
Percent_Rank
FROM Students;
```

|   | A-Z student ▼ | 123 score ▼ | 123 percent_rank ▼ |
|---|---------------|-------------|--------------------|
| 1 | Elena         | 92          | 0                  |
| 2 | Jose          | 90          | 0.25               |
| 3 | Pepe          | 85          | 0.5                |
| 4 | Marco         | 85          | 0.5                |
| 5 | Manolo        | 78          | 1                  |

- **cume\_dist**: cummulative distribution: It displays the number of values in the set preceding and including in the specified order divided by the number of rows.

```
SELECT Student, Score,
CUME_DIST() OVER (ORDER BY Score
DESC) AS Cume_Dist
FROM Students;
```

|   | A-Z student ▼ | 123 score ▼ | 123 cume_dist ▼ |
|---|---------------|-------------|-----------------|
| 1 | Elena         | 92          | 0.2             |
| 2 | Jose          | 90          | 0.4             |
| 3 | Pepe          | 85          | 0.8             |
| 4 | Marco         | 85          | 0.8             |
| 5 | Manolo        | 78          | 1               |

## Other rank functions

- **row\_number:**

```
SELECT Student, Score,
ROW_NUMBER() OVER (ORDER BY Score DESC) AS Row_Number
FROM Students;
```

|   | A-Z student | 123 score | 123 row_number |
|---|-------------|-----------|----------------|
| 1 | Elena       | 92        | 1              |
| 2 | Jose        | 90        | 2              |
| 3 | Pepe        | 85        | 3              |
| 4 | Marco       | 85        | 4              |
| 5 | Manolo      | 78        | 5              |

- **ntile(x):** cuantile

- Divides the rows in the partition in x buckets with the same number of rows

```
SELECT Student, Score,
NTILE(2) OVER (ORDER BY Score
DESC) AS ntile
FROM Students;
```

|   | A-Z student | 123 score | 123 ntile |
|---|-------------|-----------|-----------|
| 1 | Elena       | 92        | 1         |
| 2 | Jose        | 90        | 1         |
| 3 | Pepe        | 85        | 1         |
| 4 | Marco       | 85        | 2         |
| 5 | Manolo      | 78        | 2         |

- Numeric functions (exp, cos, ln, ...) **SELECT EXP(2);**
- Aggregated (std, var, corr, regr, ...) **SELECT STDDEV(Score) FROM Students;**

|   | A-Z student | 123 score | 123 previous_score | 123 next_score |
|---|-------------|-----------|--------------------|----------------|
| 1 | Elena       | 92        | [NULL]             | 90             |
| 2 | Jose        | 90        | 92                 | 85             |
| 3 | Pepe        | 85        | 90                 | 85             |
| 4 | Marco       | 85        | 85                 | 78             |
| 5 | Manolo      | 78        | 85                 | [NULL]         |

- Frame functions: lag, lead, ...

```
SELECT Student,Score,
LAG(Score) OVER (ORDER BY Score DESC) AS Previous_Score,
LEAD(Score) OVER (ORDER BY Score DESC) AS Next_Score
FROM Students;
```

SQL:1999 allows the use of **nulls first** and **nulls last**. It serves to define if nulls appear before or after non-null values in the sort ordering. By default, NULLS FIRST for DESC order and NULLS LAST for ASC order.

```
SELECT Student,Score,ROW_NUMBER() OVER (ORDER BY Score ASC) AS Default_Order,ROW_NUMBER()
OVER (ORDER BY Score ASC NULLS FIRST) AS
Nulls_First_Order,ROW_NUMBER() OVER (ORDER BY
Score ASC NULLS LAST) AS Nulls_Last_Order
FROM Students;
```

|   | A-Z student | 123 score | 123 default_order | 123 nulls_first_order | 123 nulls_last_order |
|---|-------------|-----------|-------------------|-----------------------|----------------------|
| 1 | Manolo      | 78        | 1                 | 2                     | 1                    |
| 2 | Pepe        | 85        | 2                 | 3                     | 2                    |
| 3 | Marco       | 85        | 3                 | 4                     | 3                    |
| 4 | Jose        | 90        | 4                 | 5                     | 4                    |
| 5 | Elena       | 92        | 5                 | 6                     | 5                    |
| 6 | Pat         | [NULL]    | 6                 | 1                     | 6                    |



**1 Multidimensional view of data**

**2 Transparency to support (ROLAP, MOLAP)** interfaz simple para el usuario, le da igual lo que hay detrás

**3 Accessibility** el usuario debe acceder de forma simple a los datos, da igual donde estén

**4 Coherent performance in reporting** resultados rápidos y consistentes

**5 Client-Server Architecture**

**6 Generic operations regarding the number of dimensions**

**7 Dynamic sparse matrix**

**8 Multiuser support**

**9 Flexibility in the definition of the dimensions: constraints, aggregations and hierarchies among them.**

**10 Intuitive handling of operators: drill, roll, slice-&-dice, pivot.**

**11 flexible report generation**

**12 No limit dimensions**