

Overview of Text Data Access



Tecnologías de Gestión de
Información No Estructurada
Prof. Dr. David E. Losada

CiTiUS

Centro Singular de Investigación
en Tecnologías Inteligentes

USC
UNIVERSIDADE
DE SANTIAGO
DE COMPOSTELA

CAMPUS VIDA
CAMPUS DE EXCELÊNCIA INTERDISCIPLINAR

Máster Interuniversitario en Tecnologías de Análisis de Datos Masivos: Big Data


Text Data Access

foundation for **text analysis**



enables **retrieval** of the most **relevant** text data to a particular problem

enables **interpretation** of any analysis, results or discovered knowledge in appropriate **context** and provides data **provenance** (origin)



pull: the **users** take the **initiative** to fetch relevant information out from the system



push: the **system** takes the **initiative** to offer relevant information to users

selecting **relevant text data** from a large collection is the **basic task of text access**

generally based on a specification of the **information need** of an analyst (a user), and can be pull and push

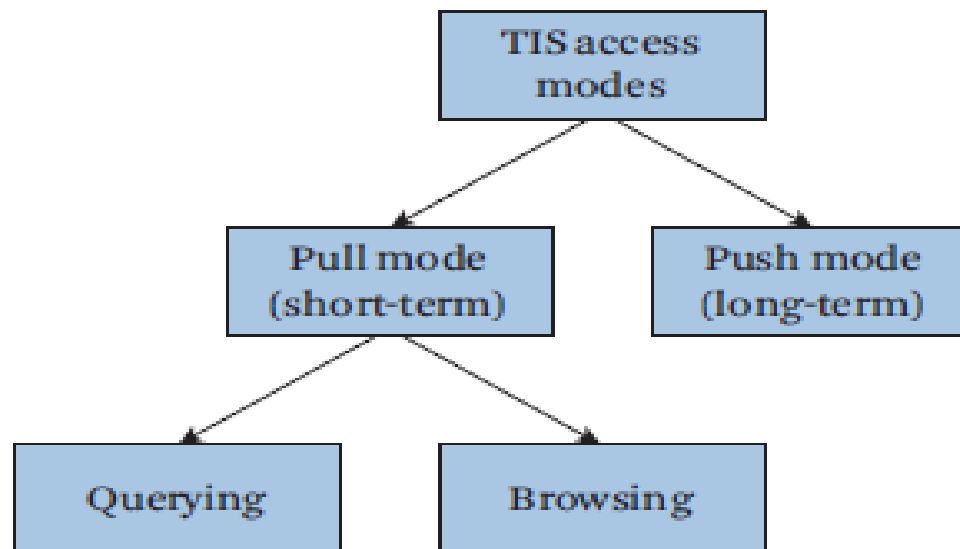


Figure 5.1 The dichotomy of text information access modes.

pull mode

the user initiates the access process to find the relevant text data, typically by using a **search engine**.

essential when a user has an ad hoc information need, i.e., a temporary information need that might disappear once the need is satisfied.



e.g., a user wants to buy a product and is interested in retrieving all the relevant opinions; after the user has purchased the product, the user would generally no longer need such information





pull mode

browsing

can be very useful when a user does not know how to formulate an effective query

or finds it inconvenient to enter a keyword query (e.g., through a smartphone), or simply wants to explore a topic with no fixed goal

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts, holding a rolled-up map and carrying a large orange backpack with a camera attached.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts, holding a rolled-up map and carrying a large orange backpack with a camera attached.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts. He is holding a rolled-up map and has a large orange backpack with a camera attached to it. He is standing on a small shadow.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts. He is holding a rolled-up map and has a large orange backpack with a camera attached to it. He is standing on a small shadow.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts, holding a rolled-up map and carrying a large orange backpack with a camera attached.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts. He is holding a rolled-up map and has a large orange backpack with a camera attached to it. He is standing on a small shadow.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts. He is holding a rolled-up map and has a large orange backpack with a camera attached to it. He is standing on a small shadow.

A cartoon illustration of a man with a large nose, wearing a yellow hat, an orange shirt, and brown shorts. He is holding a rolled-up map and has a large orange backpack with a camera attached to it. He is standing on a small shadow.

push mode

the system initiates the process to **recommend** a set of relevant information **items** to the user.

generally more useful to satisfy **long-standing information needs**

e.g., a researcher's research interests can be relatively stable over time.

the information **stream** (i.e., published research articles) is **dynamic**

although a user can regularly search for relevant literature, it is more desirable for a recommender (also called filtering) system to monitor the dynamic information stream and “push” any relevant articles to the user based on the matching of the articles with the user's interests (e.g., in the form of an email).

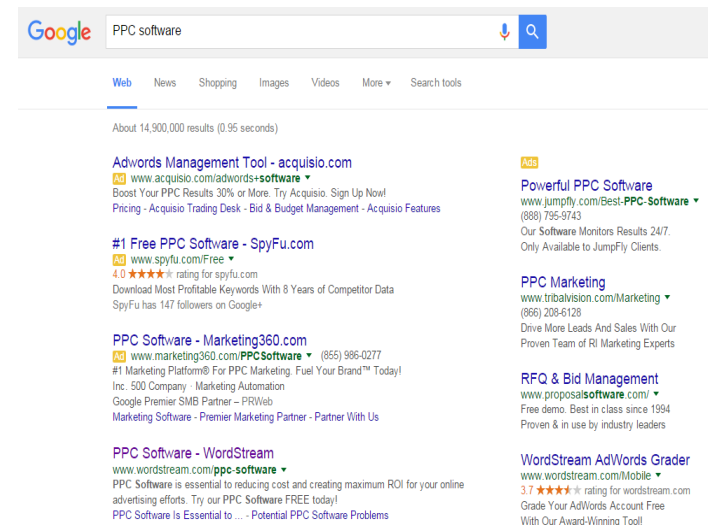


push mode

producer-initiated recommendation (or selective dissemination of information, SDI).

the producer of information has an interest in disseminating the information among relevant users, and would push an information item to such users.

e.g. advertising of products on SERP



the recommendations can be delivered through email notifications
or recommended through a SERP

short/long term information needs

short-term: most often associated with pull mode



temporary and usually satisfied through search or navigation in the information space



long-term: most often associated with push mode.

can be better satisfied through filtering or recommendation where the system would take the initiative to push the relevant information to a user.

ad-hoc retrieval extremely important

ad hoc information needs **show up far more frequently** than long-term info needs




the techniques effective for ad hoc retrieval can usually be re-used for filtering and recommendation as well.

in the case of long-term information needs, it is possible to collect and exploit **user feedback**



due to the availability of training data,
the problem of filtering/recommendation
can usually be solved by using supervised ML



in adhoc retrieval, we do not have much feedback information from a user (i.e., little training data for a particular query).

ad-hoc retrieval difficult!

ejemplo

multimode interactive access

push/pull modes integrated in the same
information access environment

querying and browsing also seamlessly integrated
to provide maximum flexibility



Figure 5.2 Sample interface of browsing with a topic map where browsing and querying are naturally integrated.



TOPIC MAP for NAVIGATION

- table (0.000313848)
- ⊖ dining (0.000139667)
 - ⊖ dining table (center)
 - round dining table (0.0104634)
 - dining table bases (0.00554701)
 - asian dining table (0.00371214)
 - door dining table (0.00301223)
 - walnut dining table (0.00291244)
 - dining table scottsdale (0.00111778)
 - dining tables (0.0554608)
 - dining chairs (0.00490976)
 - dining set (0.00191779)
 - dinning table (0.0015674)
 - patio table (0.00144361)
 - kitchen table (0.00119744)
 - buffet table (0.00056739)
 - pub table (0.000498793)
 - bar table (0.000487214)
 - coffee table (0.000159963)

Click-Through

- <http://www.diningtableguide.com/> (1)
- <http://www.diningtables.com/> (1)
- <http://www.flickr.com/photos/17425845@n00/121705244/in/pool-69453349@n00> (1)
- <http://www.glass-dining-table.net/> (1)

Search Result for "dining table"

[Dining Tables - Overstock™ Shopping - The Best Prices Online](#)

<http://www.overstock.com/Home-Garden/Dining-Tables/2021/subcat.htmlsa=Uei=hUUtVaabHMPAsAWd3IGIBgved=0CCwQFjAAusg=AFQjCNGk>

[Dining Table and Chairs Dining Tables | Pottery Barn](#)

[www.potterybarn.com/shop/furniture-upholstery/tables/CachedSimilarPotteryBarns dining ta](http://www.potterybarn.com/shop/furniture-upholstery/tables/CachedSimilarPotteryBarns%20dining%20tables)
Choose from a wide range of dining tables and chairs in classic styles and...

<http://www.potterybarn.com/shop/furniture-upholstery/tables/sa=Uei=hUUtVaabHMPAsAWd3IGIBgved=0CDcQFjABusg=AFQjCNGC1FkPZSCWB5NE>

[Dining tables - Up to 4 seats Up to 6 seats - IKEA](#)

[www.ikea.com/us/en/catalog/categories/departments/dining/21825/CachedSimilarDining table](http://www.ikea.com/us/en/catalog/categories/departments/dining/21825/CachedSimilarDining%20tables)
games, helping with homework or just lingering after a meal, they're where you share...

<http://www.ikea.com/us/en/catalog/categories/departments/dining/21825>

Figure 5.2 Sample interface of browsing with a topic map where browsing and querying are naturally integrated.

Querying (long-range jump)

a user submits a new query through
the search box
the search results from a search engine
will be shown in the right pane.

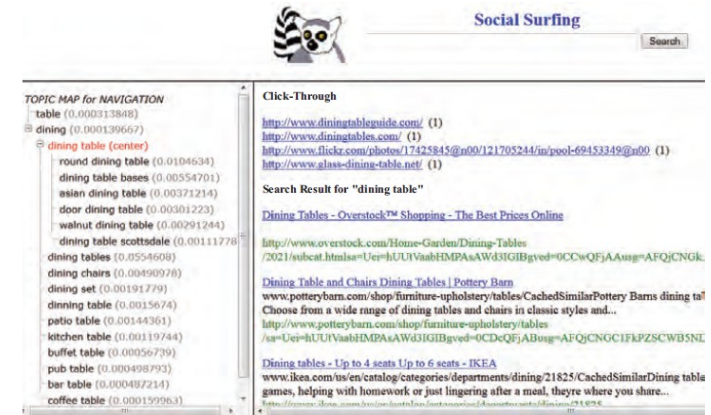


Figure 5.2 Sample interface of browsing with a topic map where browsing and querying are naturally integrated.

Navigating on the map (short-range walk)

the relevant part of a topic map is also shown on the left pane to facilitate browsing
when a user clicks on a map node, this pane will be refreshed and a local view with the clicked node as the current focus will be displayed
a user can thus zoom into a child node, zoom out to a parent node, or navigate into a horizontal neighbor node. Such a map enables the user to “walk” in the information space to browse into relevant documents without needing to reformulate queries.

Viewing a topic region



Figure 5.2 Sample interface of browsing with a topic map where browsing and querying are naturally integrated.

the user may double-click on a topic node on the map to view the documents covered in the topic region

the search result pane would be updated with new results corresponding to the documents in the selected topic region.

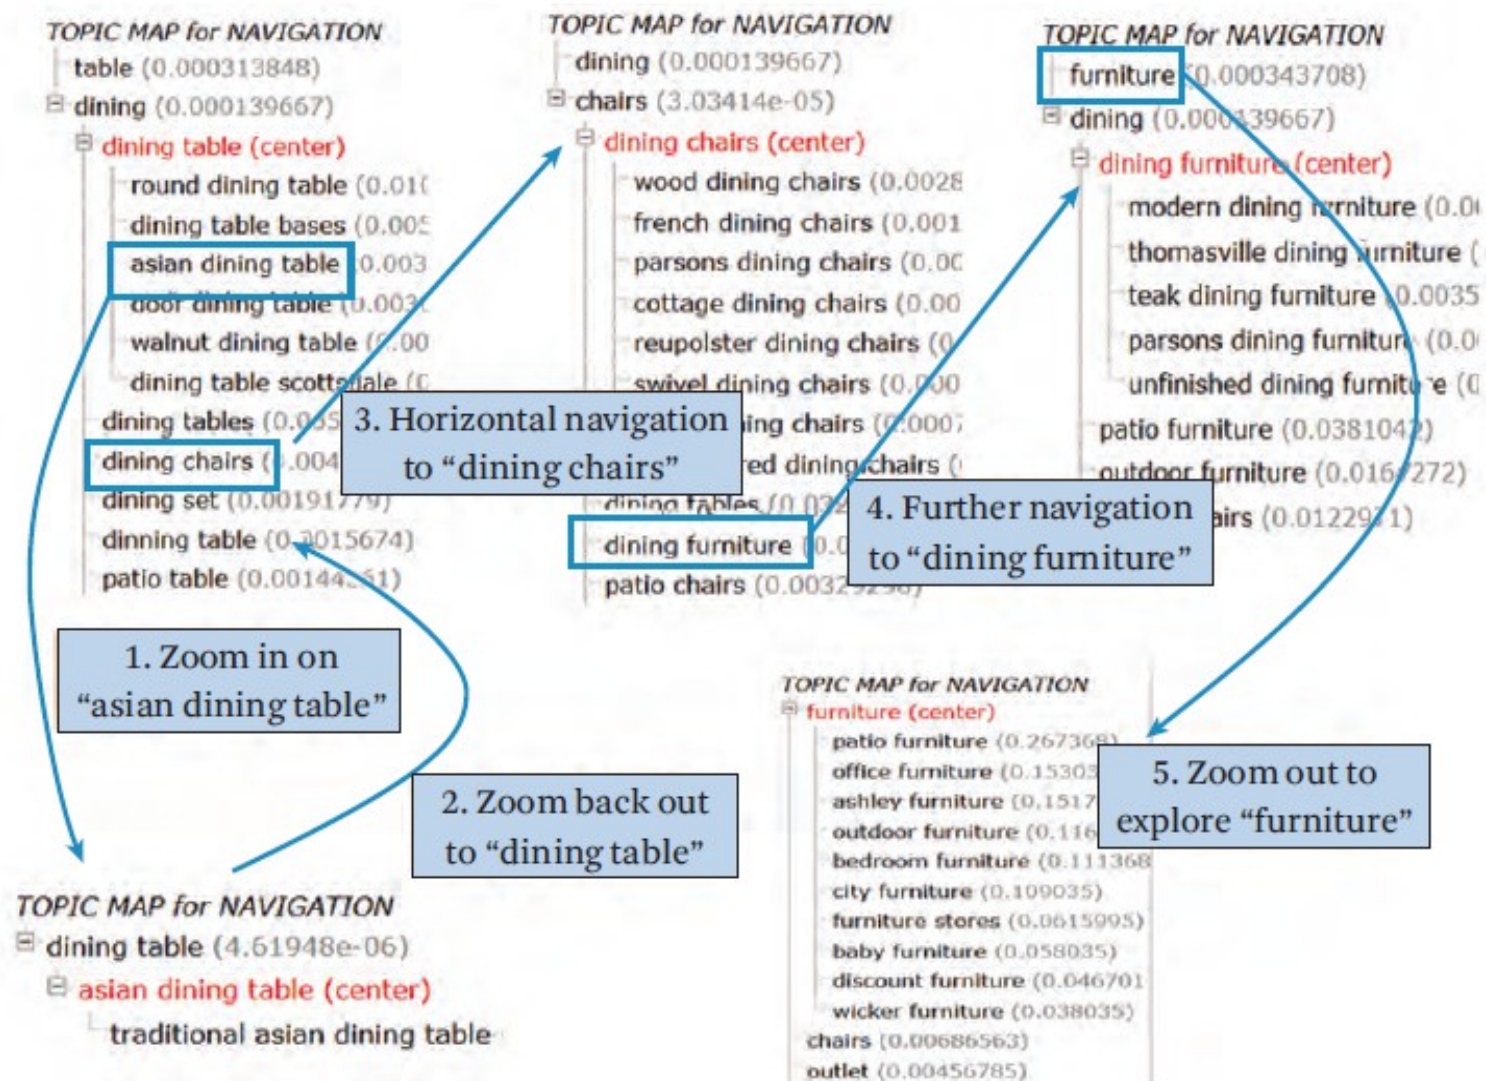


Figure 5.3 A sample trace of browsing showing how a user can navigate in the information space without querying.

aqui termina el ejemplo !!!!!

text retrieval



the most important tool for supporting **text data access**
is a **search engine**

SEs directly provide support for **querying** and can be easily extended to provide **recommendation/browsing**

the **techniques** used to implement an effective SE are often also useful for implementation of a **recommender** system as well as many **text analysis** functions



the problem of text retrieval (tr)

TR: to use a **query** to **find relevant documents** in a **collection** of text documents.

frequently needed task

(users often have temporary ad hoc information needs)



retrieval techniques for **non-textual data** are less mature (and tend to rely on TR to match a keyword query with companion text data)

For example, the current image search engines on the Web are essentially a TR system where each image is represented by a text document consisting of any associated text data with the image (e.g., title, caption,).



easy or hard

depending on **specific queries** and
specific collections




in web search, finding homepages is generally easy
but finding out people's opinions about some topic
(e.g., U.S. foreign policy) is much harder

why is text retrieval difficult?

a **query** is usually quite **short** and **incomplete**
(no formal language like SQL)

the **information need** may be difficult to describe precisely, especially when the user isn't familiar with the topic



why is text retrieval difficult?

precise **understanding** of the document **content** is difficult.



what counts as the correct answer is **subjective**


even when human experts judge the relevance of documents, they may disagree with each other

lack of clear semantic structures and difficulty in NL understanding

why is text retrieval difficult?

current SEs work very well for
navigational queries and **simple**,
popular informational queries

hay tres tipos, navegacioanl, transaccional, y otra



but in the case where a user has a **complex**
information need (e.g. analyzing opinions about
products to buy, or researching medical information
about some symptoms), they often work **poorly**

why is text retrieval difficult?

little or **no support** to help users **digest** and exploit the retrieved information

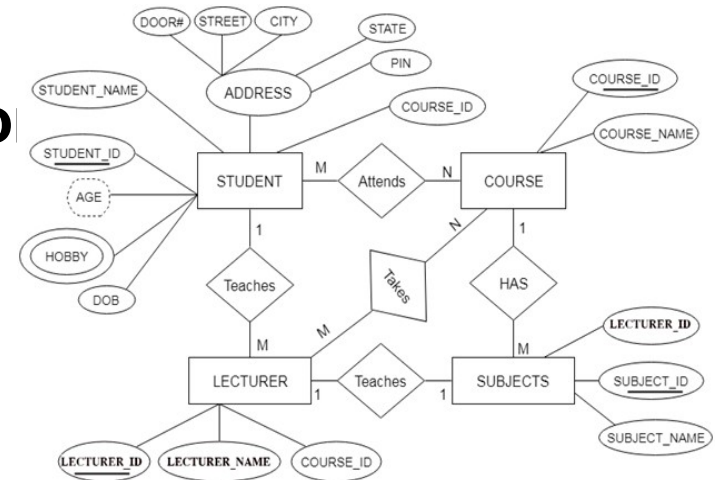


a user would still have to sift through a long list of documents and read them in detail to fully digest the knowledge buried in text data in order to perform their task at hand

text retrieval vs data retrieval

both tasks help users **find relevant info**

the **data managed** by a search engine
and a DB system are **different**



In **DBs**, the data are **structured**, each **field** has a clearly defined meaning according to a **schema**. Tables with well-specified columns.


Album	Release Date	Label	Genre
Phonology	3/1/1984	1 Rock	
Phonology	3/5/1978	1 Rock	
Colours of Passion	5/5/1980	5 Rock	
Wishes Brew	3/28/1978	4 Jazz	
Wind of Blue	8/1/1979	4 Jazz	
Goodish Stand the Weather	5/15/1984	6 Blues	
Wherever in Time	9/29/1966	1 Rock	
Notes of Mind	5/16/1983	3 Rock	
Killers	2/2/1981	3 Rock	
No Prayer for the Dying	10/1/1990	3 Rock	
Travis Road	6/13/1983	6 Blues	
Disappointed	9/28/2005	5 Hip Hop	
The Dogfather	11/12/1996	9 Hip Hop	
Let's Get to the King	8/23/2003	7 Rock	
Destiny Fulfilled	11/15/2004	8 Pop	
Back	5/12/2005	5 Hip Hop	
The Book of Luke	1/1/2005	3 Rock	
Black Ice	7/1/2006	6 Rock	
Black Ice	10/17/2008	4 Rock	
Love Songs	1/29/2012	8 Rock	

For example, in a bank DB, one field may be customer names, another may be the address, and yet another may be the balance



text retrieval vs data retrieval

the data managed by a search engine are **unstructured**
text is **difficult** for computers to **understand**



el sistema es el que tiene que
ver que es o no relevante para
cada búsqueda

even if a sentence says a person lives in a particular
address, it remains difficult for the computer to
answer a query about the address of a person in
response to a keyword query

text retrieval vs data retrieval

the **queries** that can be supported by the two are also **different**

A **DB query** clearly specifies the constraints on the fields of the data table, and thus the expected results are very **well specified** with **no ambiguity**

```
SELECT clause { SELECT  
                first_name  
FROM clause { FROM  
                employees  
WHERE clause { WHERE  
                YEAR(hire_date) = 2000
```

In a **search engine**, the queries are generally **keyword queries**, which are only a **vague specification** of what documents should be returned.

For example, in the case of searching for relevant literature to a research problem, the user is unlikely able to clearly and completely specify which documents should be returned

text retrieval vs data retrieval

DB search: we can retrieve very **specific data elements** (e.g., specific columns)

TR: we are generally only able to retrieve a set of relevant **documents**.


A SE can also retrieve **passages**, but it is difficult to retrieve specific entities or attribute values as we can in a database.



text retrieval vs data retrieval

In **DBs** there is no challenge in determining which data elements satisfy the user's query

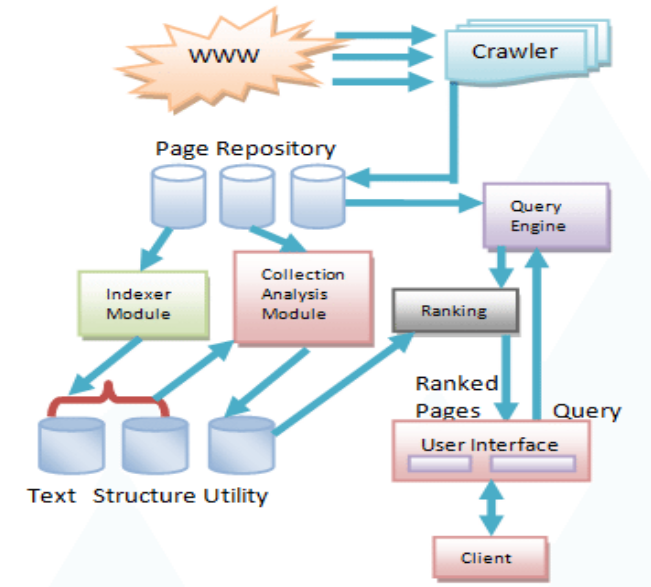
a major remaining challenge is how to find the answers **as quickly as possible**



many **queries** being **issued at the same time**

text retrieval vs data retrieval

the **efficiency** challenge also exists
in a **SE**



but a more important challenge is to figure out
which documents should be returned



text retrieval vs data retrieval

In **DBs**, it is also crucial to maintain the **integrity** of the data; e.g., to ensure no inconsistency occurs due to failures



text retrieval vs data retrieval

In TR, modeling a user's **information need** and search **tasks** is important, again due to the difficulty for a user to clearly specify information needs and the difficulty in NLP



text retrieval vs data retrieval

In TR, there is no mathematical way to prove that one answer/method is better than another

TR relies on **empirical evaluation** using some **test collections** and **users**



In DB, the main issue is **efficiency**, and one can prove one algorithm is better than another by analyzing the **computational complexity** or do some simulation study.

O (Big Oh), Ω (Omega), and Θ (Theta)

text retrieval vs data retrieval

DBs. Traditional field. Widespread applications in virtually every domain with a well-established strong industry



IR community. interdisciplinary community involving **library** and **information science** and **computer science**
strong industry base since the early 1990s



as more and more online information is available, the search engine technologies (including components such as machine learning and natural language processing) will continue to grow



the task of text retrieval

Given a **document collection** (a set of unordered text documents), the task of TR retrieval can be defined as using a **user query** (i.e., a description of the user's **information need**) to identify a subset of documents that can **satisfy the user's information need**



the task of text retrieval



$V = \{w_1, \dots, w_N\}$ be a **vocabulary** (set of all the words in a particular natural language)

user's query $q = q_1, q_2, \dots, q_m$, a sequence of words, $q_i \in V$.

doc $d_i = d_{i1}, \dots, d_{im}$, a sequence of words, $d_{ij} \in V$.

In general, a **query** is much **shorter** than a document but some retrieval tasks have large queries (e.g. patent retrieval)



the task of text retrieval



text collection $C = \{d_1, \dots, d_M\}$ is a set of textual docs.

$R(q)$, subset of docs, i.e., $R(q) \subset C$, which are relevant to the user's query q

the user's query is only a “**hint**”

different users may use the same query to intend to retrieve somewhat different sets of relevant documents



the task of text retrieval



it is **unrealistic** to expect a computer to return **exactly** the set $R(q)$ (unlike the case in database search)



return an **approximation** of $R(q)$, denoted as $R'(q)$

how can a computer compute $R'(q)$?

document selection vs. document ranking

document selection

a **binary classifier** to classify a document as either **relevant or non-relevant** with respect to a particular query.

indicator function $f(q, d) \in \{0, 1\}$

$R'(q) = \{d \mid f(q, d) = 1, d \in C\}$.



the system estimates the “**absolute relevance**”

document ranking

rank documents and let the user decide a cutoff

a ranking function $f(q, d) \in \mathbb{R}$

rank all the documents in descending values of this ranking function

$R'(q)$ is defined partly by the system and partly by the user

the user implicitly chooses a score threshold θ based on the rank position where he stopped

$$R'(q) = \{d \mid f(q, d) \geq \theta\}$$

the system only needs to estimate the “**relative relevance**” of documents



document ranking



estimation of **relative relevance** is intuitively **easier** than that of absolute relevance

ranking is generally preferred to document selection

difficulty for a user to prescribe the exact criteria for selecting relevant documents, **the binary classifier is unlikely accurate**

query over-constrained, there may be **no relevant documents** matching all the query words, so forcing a binary decision may result in **no delivery of any search result**

query under-constrained (too general), **too many documents** matching the query, resulting in **over-delivery**



even if the classifier can be accurate, a user would still benefit from **prioritization** of the matched relevant documents for examination.

some relevant docs are more useful than others (**relevance degrees**)

probability ranking principle (PRP)

(Robertson 1997) the strategy of **ranking** shown to be **optimal theoretically** under two assumptions

a ranked list of documents in descending order of predicted relevance is the optimal strategy under the following two assumptions:

1. The utility of a document to a user is independent of the utility of any other document.
2. A user will browse the results sequentially.

